



10-301/10-601 Introduction to Machine Learning

Machine Learning Department
School of Computer Science
Carnegie Mellon University

Decision Trees

Matt Gormley & Geoff Gordon
Lecture 3
Sep. 3, 2025

Q&A

Q: How do these In-Class Polls work?

- A:**
- Sign into **Google Form** (click [Poll] link on Schedule page <http://mlcourse.org/schedule.html>) using **Andrew Email**
 - Answer **during lecture for full credit**, or within 24 hours for half credit
 - Avoid the **toxic option** which gives negative points!
 - 8 “free poll points” but can’t use more than 3 free polls consecutively. All the questions for one lecture are worth 1 point total.

Latest Poll link: <http://poll.mlcourse.org>

broken

First In-Class Poll

Question: Which of the following did you bring to class today? Select all that apply.

A. Smartphone



B. Flip phone



C. Pay phone



D. No phone



E. None of the above



Answer:

TOXIC →



→ do not pick


Reminders

- **Homework 1: Background**
 - Out: Mon, Aug 25
 - Due: Wed, Sep 3 at 11:59pm
 - unique policy for this assignment: we will grant (essentially) any and all extension requests
- **Homework 2: Decision Trees**
 - Out: Wed, Sep. 3
 - Due: Mon, Sep. 15 at 11:59pm

MAKING PREDICTIONS WITH A DECISION TREES

Background: Recursion

- Def: a **binary search tree** (BST) consists of nodes, where each node:
 - has a value, v
 - up to 2 children
 - all its left descendants have values less than v , and its right descendants have values greater than v
- We like BSTs because they permit search in $O(\log(n))$ time, assuming n nodes in the tree



Node Data Structure

```
class Node:  
    int value  
    Node left  
    Node right
```

contains (root, 6) = true

Iterative Search

```
def contains(node, key):  
    cur = node  
    while True:  
        if key < cur.value & cur.left != null:  
            cur = cur.left  
        else if cur.value < key & cur.right != null:  
            cur = cur.right  
        else:  
            break  
    return key == cur.value
```

contains (root, 6) →
contains (Node(3), 6)
→ contains (Node(6), 6)
→ return true

Recursive Search

```
def contains(node, key):  
    → if key < node.value & node.left != null:  
        return contains(node.left, key)  
    → else if node.value < key & node.right != null:  
        return contains(node.right, key)  
    else:  
        return key == node.value
```

Algorithms for Classification

Algorithm 3 decision stump: based on a single feature, x_d , predict the most common label in the training dataset among all data points that have the same value for x_d

predictions	y	x_1	x_2	x_3	x_4
	allergic?	hives?	sneezing?	red eye?	has cat?
-	-	Y	N	N	N
+	-	N	Y	N	N
+	+	Y	Y	N	N
-	-	Y	N	Y	Y
+	+	N	Y	Y	N


But why use
just one
feature...

Nonzero training error, but
perhaps still better than
the memorizer

Example
decision stump:
$$h(x) = \begin{cases} + & \text{if sneezing} = Y \\ - & \text{otherwise} \end{cases}$$

From Decision Stump to Decision Tree


predictions	y	x ₁	x ₂	x ₃	x ₄
	allergic?	hives?	sneezing?	red eye?	has cat?
-	-	Y	N	N	N
-	-	N	Y	N	N
+	+	Y	Y	N	N
-	-	Y	N	Y	Y
+	+	N	Y	Y	N



From Decision Stump to Decision Tree

~~✗~~

	y	x_1	x_2	x_3	x_4
predictions	allergic?	hives?	sneezing?	red eye?	has cat?
-	-	Y	N	N	N
-	-	N	Y	N	N
+	+	Y	Y	N	N
-	-	Y	N	Y	Y
+	+	N	Y	Y	N



Decision Tree: In-Class Activity

1. Group 1: Answer the questions to determine which leaf node corresponds to your feature values
2. Group 2: (part 1)
 - a) Take a blue sticky note if you prefer dogs to cats; otherwise, take a red sticky note
 - b) Answer the questions to determine which leaf node corresponds to your feature values and place your sticky note there
3. Group 2: (part 2)
 - a) Answer the new question to determine which new leaf node to move your sticky note to

Decision Tree: Prediction

```
def h(x'):
```

Let current node = root

```
while(true):
```

if current node is internal (non-leaf): ↪

 Let m = attribute associated with current node

 Go down branch labeled with value x'_m

if current node is a leaf:

 return label y stored at that leaf

branches = { "cold" : Node(), "hot" : Node() }

```
class Node:
```

 str type // "leaf" or "internal"

 y vote // label for leaf node

 {} branches // map from feature


 // values to Node objects

 int m // feature for internal node

Decision Tree: Prediction

Algorithm 4 decision tree: recursively walk from root to a leaf, following the attribute values labeled on the branches, and return the label at the leaf

		y	x_1	x_2	x_3	x_4
		allergic?	hives?	sneezing?	red eye?	has cat?
predictions	-	-	Y	N	N	N
-	-	N	Y	Y	N	N
+	+	Y	Y	Y	N	N
-	-	Y	N	Y	Y	Y
+	+	N	Y	N	Y	N



Zero training error!

Decision Tree: Prediction (Iterative)

```
def h(x'):  
    Let current node = root  
    while(true):  
        if current node is internal (non-leaf):  
            Let m = attribute associated with current node  
            Go down branch labeled with value  $x'_m$   
        if current node is a leaf:  
            return label y stored at that leaf
```

Question: The original $h(x')$ pseudocode is an iterative implementation. Can you implement $h(x')$ recursively?

```
class Node:  
    str type // “leaf” or “internal”  
    Y vote // label for leaf node  
    {} branches // map from feature  
                // values to Node objects  
    int m // feature for internal node
```

Decision Tree: Prediction (Recursive)

```
def h( $\vec{x}$ ):  
    return h_recurse(root,  $\vec{x}$ )  
  
def h_recurse(node,  $\vec{x}$ ):  
    if node.type == "leaf":  
        return node.vote  
    else:  
        m = node.m  
        next = branches [ $x_m$ ]  
        return h_recurse(next,  $\vec{x}$ )
```

Question: The original $h(x')$ pseudocode is an iterative implementation. Can you implement $h(x')$ recursively?

```
class Node:  
    str type // "leaf" or "internal"  
    Y vote // label for leaf node  
    {} branches // map from feature  
                // values to Node objects  
    int m // feature for internal node
```


Decision Tree Example

Learned from medical records of 1000 women (Sims et al., 2000)

Negative examples are C-sections

[833+, 167-] .83+ .17-

```
Fetal_Presentation = 1: [822+, 116-] .88+ .12-
| Previous_Csection = 0: [767+, 81-] .90+ .10-
| | Primiparous = 0: [399+, 13-] .97+ .03-
| | Primiparous = 1: [368+, 68-] .84+ .16-
| | | Fetal_Distress = 0: [334+, 47-] .88+ .12-
| | | Fetal_Distress = 1: [34+, 21-] .62+ .38-
| | Previous_Csection = 1: [55+, 35-] .61+ .39-
Fetal_Presentation = 2: [3+, 29-] .11+ .89-
Fetal_Presentation = 3: [8+, 22-] .27+ .73-
```



LEARNING A DECISION TREE

Decision Tree Learning

- *Definition:* a **splitting criterion** is a function that measures the effectiveness of splitting on a particular attribute
- Our decision tree learner **selects the “best” attribute** as the one that maximizes the splitting criterion
- Lots of options for a splitting criterion:
 - error rate (minimize)
 - accuracy = 1 - error rate (maximize)
 - Mutual information (maximize)
 - Gini gain (maximize)

Decision Tree Learning

```
def train(D):  
    root = train_recurse(D)  
    store root
```

```
def train_recurse(D'):  
    Let p = new Node()  
  
    [Base Case] IF (a) D' is empty  
    (b) D' has all labels identical  
    (c) for each feature, all values in D' are identical  
    (d) depth of node ≥ max_depth  
        p.type = "leaf"  
        p.vote = majority_vote(D')  
        return p
```

```
[Recursive Case] Else:  
    p.type = "internal"  
    p.m = best feature according to splitting criterion on D'  
    = argmaxm ∈ {1, ..., M} splitting_criterion(D', m)
```

for each value v of feature p.m:

$D_{x_m=v} = \{(x, y) \in D' : x_m = v\}$ ← partition of D'

child_v = train_recurse($D_{x_m=v}$) ← recursion


p.branches{v} = child_v ← add a branch w/ label v
return p

Decision Tree Learning Example


Dataset:

	A	B	C
y	Y	N	N
-	Y	N	Y
-	Y	N	N
+	N	N	Y
+	Y	Y	N
+	Y	Y	Y
+	Y	Y	N
+	Y	Y	Y

at root D has $\{5+, 3-\}$



In-Class Exercise: Using error rate as the splitting criterion, what decision tree would be learned?



Recursive Training for Decision Trees

- def train(dataset D):
 - Let p = new Node()
 - **Base Case:** If (1) all labels $y^{(i)}$ in D' are identical (2) D' is empty (3) for each attribute, all values are identical
 - p.type = Leaf // The node p is a leaf node
 - p.label = majority_vote(D') // Store the label
 - return p
 - **Recursive Step:** Otherwise
 - Make an internal node
 - p.type = Internal // The node p is an internal node
 - Pick the *best* attribute X_m according to splitting criterion
 - p.attr = argmax_m splitting_criterion(D', X_m) // Store the attribute on which to split
 - For each value v of attribute X_m :
 - $D_{X_m=v} = \{(x,y) \text{ in } D' : x_m = v\}$ // Select a partition of the data
 - $\text{child}_v = \text{train}(D_{X_m=v})$ // Recursively build the child
 - p.branches[v] = child_v // Create a branch with label v
 - return p

SPLITTING CRITERION: ERROR RATE

Decision Tree Learning Example

Dataset:

Label Y, Features A and B

Y	A	B
-	1	0
-	1	0
+	1	0
+	1	0
+	1	1
+	1	1
+	1	1
+	1	1

Poll Question 2

Which attribute would **error rate** select for the next split?

1. A 12%
2. B 8%
3. A or B (tie) 79%
4. Neither ← toxic

Decision Tree Learning Example

Dataset:

Label Y, Features A and B

Y	A	B
-	1	0
-	1	0
+	1	0
+	1	0
+	1	1
+	1	1
+	1	1
+	1	1




Decision Tree Learning Example

Dataset:

Label Y, Features A and B

Y	A	B
-	1	0
-	1	0
+	1	0
+	1	0
+	1	1
+	1	1
+	1	1
+	1	1



Error Rate

$$\text{error}(h_A, D) = 2/8$$

$$\text{error}(h_B, D) = 2/8$$

error rate treats
attributes A and B as
equally good

Decision Tree Learning

- *Definition:* a **splitting criterion** is a function that measures the effectiveness of splitting on a particular attribute
- Our decision tree learner **selects the “best” attribute** as the one that maximizes the splitting criterion
- Lots of options for a splitting criterion:
 - error rate (minimize)
 - accuracy = 1 - error rate (maximize)
 - Mutual information (maximize)
 - Gini gain (maximize)

SPLITTING CRITERION: MUTUAL INFORMATION

R.V. over say
labels

Entropy

- The **entropy** of a *random variable* describes the uncertainty of its outcome: the higher the entropy, the less certain we are about what the outcome will be.

$$H(X) = - \sum_{v \in V(X)} P(X = v) \log_2(P(X = v))$$

where X is a (discrete) random variable

$V(X)$ is the set of possible values X can take on

Entropy

S \leftrightarrow \hat{S} Labels
set \hat{S} labels

- The **entropy** of a set describes how uniform or pure it is: the higher the entropy, the more impure or “mixed-up” the set is

$$H(S) = - \sum_{v \in V(S)} \frac{|S_v|}{|S|} \log_2 \left(\frac{|S_v|}{|S|} \right)$$

where S is a collection of values, $V(S)$ is the set of unique values in S , S_v is the collection of elements in S with value v , $|S_v|$ is the number of elements in S_v , and $|S|$ is the number of elements in S .

$V(S)$ is the set of unique values in S

S_v is the collection of elements in S with value v
two possible values

- If all the elements in S are the same, then

$$H(S) = - \left(\frac{9}{9} \log_2 \frac{9}{9} + \frac{0}{9} \log_2 \frac{0}{9} \right) = 0$$

0 + 0

Entropy

- The **entropy** of a *set* describes how uniform or pure it is: the higher the entropy, the more impure or “mixed-up” the set is

$$H(S) = - \sum_{v \in V(S)} \frac{|S_v|}{|S|} \log_2 \left(\frac{|S_v|}{|S|} \right)$$

where S is a collection of values,

$V(S)$ is the set of unique values in S

S_v is the collection of elements in S with value v

- If S is split fifty-fifty between two values, then

$$H(S) = - \left(\frac{5}{10} \log \frac{5}{10} + \frac{5}{10} \log \frac{5}{10} \right) = 1$$

Mutual Information

Y *label*
 X *feature candidate*

- The **mutual information** between *two random variables* describes how much clarity knowing the value of one random variables provides about the other

$$\begin{aligned} I(Y; X) &= \underbrace{H(Y)}_{\text{---}} - \underbrace{H(Y|X)}_{\text{---}} \\ &= H(Y) - \sum_{v \in V(X)} P(X = v)H(Y|X = v) \end{aligned}$$

where X and Y are random variables

$V(X)$ is the set of possible values X can take on

$H(Y|X = v)$ is the conditional entropy of Y given $X = v$

Mutual Information

- The **mutual information** between a feature and the label describes how much clarity knowing the feature provides about the label

$$I(y; x_d) = H(y) - H(y|x_d)$$

weighted avg of entropies at children

entropy at parent $\quad = \overbrace{H(y)}^{\text{parent}} - \sum_{v \in V(x_d)} f_v * \underbrace{H(Y_{x_d=v})}_{\text{entropy of each child}}$

where x_d is a feature and y is the set of all labels

$V(x_d)$ is the set of possible values x_d can take on

f_v is the fraction of data points where $x_d = v$

$Y_{x_d=v}$ is the set of all labels where $x_d = v$

Decision Tree Learning Example

Dataset:

Label Y, Features A and B

Y	A	B
-	1	0
-	1	0
+	1	0
+	1	0
+	1	1
+	1	1
+	1	1
+	1	1

Poll Question 3

SKIP

Which feature would
mutual information
select for the next
split?

1. A
2. B
3. A or B (tie)
4. Neither

Decision Tree Learning Example

Dataset:

Label Y, Features A and B

Y	A	B
-	1	0
-	1	0
+	1	0
+	1	0
+	1	1
+	1	1
+	1	1
+	1	1


=\

Decision Tree Learning Example

Dataset:

Label Y, Features A and B

Y	A	B
-	1	0
-	1	0
+	1	0
+	1	0
+	1	1
+	1	1
+	1	1
+	1	1



Mutual Information

$$H(Y) = -2/8 \log(2/8) - 6/8 \log(6/8)$$

$$H(Y|A=0) = \text{"undefined"}$$

$$H(Y|A=1) = -2/8 \log(2/8) - 6/8 \log(6/8) \\ = H(Y)$$

$$H(Y|A) = P(A=0)H(Y|A=0) + P(A=1)H(Y|A=1) \\ = 0 + H(Y|A=1) = H(Y)$$

$$I(Y; A) = H(Y) - H(Y|A=1) = 0$$

$$H(Y|B=0) = -2/4 \log(2/4) - 2/4 \log(2/4)$$

$$H(Y|B=1) = -0 \log(0) - 1 \log(1) = 0$$

$$H(Y|B) = 4/8(0) + 4/8(H(Y|B=0))$$

$$I(Y; B) = H(Y) - 4/8 H(Y|B=0) > 0$$